

# Operating Systems

Sateesh K. Peddoju


1

1

## Process Synchronization

### Dining Philosopher's Problem

- 5 Philosophers
- 5 chopsticks (forks)
- Random times eat
- Must pick
  - One right & One left
- Issues
  - Deadlock or Starvation
- Solutions?
  - Only 4 to be hungry
  - Allow – if both sticks are available
  -



2

2

## Process Synchronization

### Barber Shop Problem

- sleeping Barber Problem
- Single Barber
- A barber's chair
- n – chairs for waiting
- Barber spends his life in shop
- Sleeps until a customer wakes up.
- Customer waits, if busy
- No chairs, leaves & comes back
- Solution
  - Mutual Exclusion

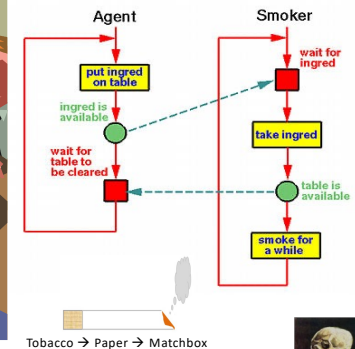


3

3

## Process Synchronization

### Cigarette Smokers Problem



Tobacco → Paper → Matchbox



Cigarette Smoking is injurious to health



4

## Process synchronization

- Producer-Consumer Problem
  - ( or Bounded-Buffer Problem)
- Readers-Writers Problem

5

## CLASSICAL PROBLEMS OF SYNCHRONIZATION

### PRODUCER – CONSUMER PROBLEM ( BOUNDED-BUFFER ) ~~WITHOUT SEMAPHORES~~

6

## Bounded-Buffer ~~without~~ semaphores

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...                /* int data; */
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

7

## Bounded-Buffer ~~without~~ semaphores

- **Producer process**

```
item nextProduced;
while (1) {
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

- **Consumer process**

```
item nextConsumed;
while (1) {
    while (counter == 0)
        ; /* do nothing */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
}
```

8

## Bounded-Buffer ~~without~~ semaphores

- The statements

```
counter++;
counter--;
```

must be performed *atomically*.

- Atomic operation means an operation that completes in its entirety without interruption.

9

9

## Bounded-Buffer ~~without~~ semaphores

- The statement “**count++**” may be implemented in machine language as:

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- The statement “**count--**” may be implemented as:

```
register2 = counter
register2 = register2 - 1
counter = register2
```

10

10

## Bounded-Buffer ~~without~~ semaphores

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.


11

11

## Bounded-Buffer ~~without~~ semaphores

- Assume **counter** is initially 5. One interleaving of statements is:

```
producer: register1 = counter (register1 = 5)
producer: register1 = register1 + 1 (register1 = 6)
consumer: register2 = counter (register2 = 5)
consumer: register2 = register2 - 1 (register2 = 4)
producer: counter = register1 (counter = 6)
consumer: counter = register2 (counter = 4)
```



- The value of **count** may be either 4 or 6, where the correct result should be 5.

12

12

## Concurrency


- Benefits
  - Communication among the processes
  - Sharing Resources
  - **Synchronization of multiple processes**
  - Allocation of processors time
- Difficulties
  - Sharing global resources (ex: global variable)
    - **Race Condition**
  - Management of allocation of resources (optimality?)
  - Programming errors difficult to locate
  - **Inconsistent data**

13

13

## Concurrency – Race Condition

- A race condition occurs when multiple processes or threads read and write data items so that the **final result depends on the order of execution of instructions** in the multiple processes



14

14

## Concurrency – OS Design Concerns

- **Multiprogramming**
  - Management of multiple processes in Uni-processor
- **Multiprocessing**
  - Management of multiple processes in Multi-processors
- **Distributed Processing**
  - Management of multiple processes executing on multiple, distributed machines

15

15

## OS Design Concerns

- Keep track of active processes
- Allocate and deallocate resources
  - Processor time
  - Memory
  - Files
  - I/O devices
- Protect data and resources
- *Result of process must be independent of the speed of execution of other concurrent processes*

16

16

## Contexts of Concurrency

- Non-concurrent vs. Concurrent code
- Structured Applications
  - Concurrency is part of programming
  - User Threads
- Operating System Structure
  - Concurrency is part of OS
  - Kernel Threads
- Requirement: Identify concurrent sections of the code.

17

17

## Critical Section

- A **section of code** within a process that requires access to **shared resources** and that **must not be executed** while another process is in a corresponding section of code.

18

18

## Critical Section Problem

- **Design a protocol** that the processes can use to cooperate.
- Each process must request to enter critical section
- Entry Section
- Exit Section
- Remainder Section

```
do {
    entry section
    critical section
    exit section
    remainder section
} while (TRUE);
```

19

19

## Solution to Critical Section Problem

### Mutual Exclusion

1. Only one process at a time is allowed in the critical section for a resource

### Progress

2. A process that is in its remainder section must do so without interfering with other processes
3. A process must not be delayed access to a critical section when there is no other process using it
4. It must not be possible for a process to be delayed indefinitely
5. A process remains inside its critical section for a finite time only
6. We assume that each process is executing at non-zero speed. However, we can make no assumption concerning the relative speed of the processes.

### Bounded Waiting

7. If another process is waiting to enter the critical section, there exists a bound or limit on the number of times current process to enter its critical section

20

20

## Mutual Exclusion

- The **requirement** that when **one** process is **in** a **critical section** that accesses shared resources, **no other** process may be **in** a **critical section** that accesses any of those shared resources.

```

/* PROCESS 1 */
void P1
{
    while (true) {
        /* preceding code */
        entercritical (Ra);
        /* critical section */
        exitcritical (Ra);
        /* following code */
    }
}

/* PROCESS 2 */
void P2
{
    while (true) {
        /* preceding code */
        entercritical (Ra);
        /* critical section */
        exitcritical (Ra);
        /* following code */
    }
}

...

/* PROCESS n */
void Pn
{
    while (true) {
        /* preceding code */
        entercritical (Ra);
        /* critical section */
        exitcritical (Ra);
        /* following code */
    }
}

```

21

21

## Meeting the Requirements...

- Software Approach:**
  - Leave the responsibility with the processes to handle the situation (neither OS nor PL will support)
  - Prone to high processing overhead and bugs
- Hardware Approach:**
  - Using special purpose machine instructions (at ISA)
  - Reduces overhead but not the general-purpose solution
- OS and PL Approach:**
  - Provide support within OS and PL itself.

22

22

## Software Approaches

- Peterson's Solution**
  - Restricted to **two processes only** (P0, P1) (Pi, Pj).
- `int turn;`  
/\* whose turn it is to enter critical section \*/
- `boolean flag[2];`  
/\* which process is ready to enter critical section \*/
- Mutual exclusion?
- Progress ?
- Bounded Wait ?

```

do {
    flag[i] = TRUE;
    turn = j;
    while (flag[j] && turn == j);

    critical section
    flag[i] = FALSE;

    remainder section
} while (TRUE);

```

23

23

## Software Approaches

- A general Solution - Lock**

```

do {
    do {
        entry section
        critical section
        exit section
        remainder section
    } while (TRUE);
} while (TRUE);

do {
    acquire lock
    critical section
    release lock
    remainder section
} while (TRUE);

```

24

24

## Software Approaches


- Disadvantage
  - Prone to high processing overhead and bugs

25

25

## Hardware Approaches

- **Interrupt Disabling**
  - uniprocessor
  - Disabling interrupts guarantees mutual exclusion
- Disadvantages:
  - Processor is limited in its ability to interleave programs
  - Multiprocessing
    - disabling interrupts on one processor will not guarantee mutual exclusion



26

## Hardware Approaches

- **Special Machine Instructions**
  - Performed in a single instruction cycle
  - Not subject to interference from other instructions
  - TestAndSet() and Swap() (Reading)

27

27

## Hardware Approaches

- Advantages
  - Applicable to any number of processes on either a single processor or multiple processors sharing main memory
  - It is simple and therefore easy to verify
  - It can be used to support multiple critical sections

28

28

## Hardware Approaches

- Disadvantages
  - Busy-waiting
    - consumes processor time
  - Starvation
    - more than one process is waiting → some process may not be selected → waiting indefinitely.
  - Deadlock
    - If a low priority process (say P1) executing in the critical region and a higher priority process (P2) needs, the higher priority process will obtain the processor.
    - If P2 tries to access a resource which P1 holds, it will be denied due to mutual exclusion.
    - P1 to wait for the critical region, P2 to wait for resource

29

29

## OS and PL Approaches

- Semaphores
  - Binary Semaphores
    - Mutex
  - Condition Variable
- Monitor
  - Event Flags
- Mailboxes/Messages
- Spinlocks

30

30

## OS and PL Approaches

### • Semaphore

```
do {
    waiting(mutex);

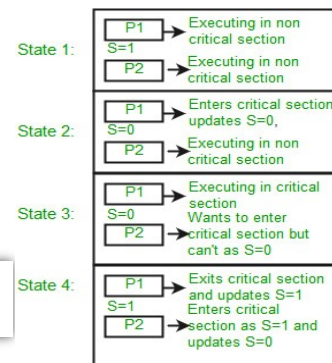
    // critical section

    signal(mutex);

    // remainder section
}while (TRUE);
```

```
wait(S) {
    while S <= 0
        ; // no-op
    S--;
```

```
signal(S) {
    S++;
```



31

31

## Semaphore Implementation

- When a process must wait on a semaphore, it is added to the list of processes
- A Signal() operation removes one process from the list of waiting processes and awakens that process.

```
typedef struct {
    int value;
    struct process *list;
} semaphore;
```

32

32



## OS and PL Approaches

### • Semaphore – Disadvantages

– Busy waiting

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

### Solution

OS should have system calls:

block()  
wakeup()

May be negative also

If negative → s = no. of processes  
waiting for s

```
signal(S) {
    S++;
}
```

```
wait(S) {
    while S <= 0
        ; // no-op
    S--;
}
```